

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
from plotly.subplots import make_subplots

from sklearn.metrics import classification_report

import warnings
warnings.filterwarnings("ignore")
```

```
DF=pd.read_csv('/content/heart.csv')
DF.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	52	1	0	125	212	0	1	168	0	1.0	2	2	
1	53	1	0	140	203	1	0	155	1	3.1	0	0	
2	70	1	0	145	174	0	1	125	1	2.6	0	0	
3	61	1	0	148	203	0	1	161	0	0.0	2	1	
4	62	0	0	138	294	1	1	106	0	1.9	1	3	

Next steps:

Generate code with DF

View recommended plots

```
nRow, nCol = DF.shape
print(f'There are {nRow} rows and {nCol} columns')
```

There are 1025 rows and 14 columns

```
DF.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   age         1025 non-null   int64
1   sex         1025 non-null   int64
2   cp          1025 non-null   int64
3   trestbps    1025 non-null   int64
4   chol        1025 non-null   int64
5   fbs         1025 non-null   int64
6   restecg     1025 non-null   int64
7   thalach     1025 non-null   int64
8   exang       1025 non-null   int64
9   oldpeak     1025 non-null   float64
10  slope       1025 non-null   int64
11  ca          1025 non-null   int64
12  thal        1025 non-null   int64
13  target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

```
DF.dtypes


age          int64
sex          int64
cp           int64
trestbps     int64
chol         int64
fbs          int64
restecg      int64
thalach      int64
exang        int64
oldpeak      float64
slope        int64
ca           int64
thal         int64
target       int64
dtype: object
```

```
DF.describe().round(2).style.background_gradient()
```

	age	sex	cp	trestbps	cho1	fbs	
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	54.430000	0.700000	0.940000	131.610000	246.000000	0.150000	
std	9.070000	0.460000	1.030000	17.520000	51.590000	0.360000	
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	
25%	48.000000	0.000000	0.000000	120.000000	211.000000	0.000000	
50%	56.000000	1.000000	1.000000	130.000000	240.000000	0.000000	
75%	61.000000	1.000000	2.000000	140.000000	275.000000	0.000000	
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	

```
DF_corr= DF.corr()
```


```
plt.figure(figsize = (14,5))
sns.heatmap(round(DF_corr,2),annot=True,cmap = 'mako')
```



```
Mis_features=['thal','ca','slope','exang','restecg','fbs','cp','sex']
DF[Mis_features] = DF[Mis_features].astype(object)
```

```
sns.countplot(data=DF,x='target')
```

<Axes: xlabel='target', ylabel='count'>



```
Dum_DF=pd.get_dummies(DF,columns=['thal','ca','slope','exang','restecg','fbs','cp','sex'],drop_first=True)
```

```
Dum_DF.head()
```

	age	trestbps	chol	thalach	oldpeak	target	thal_1	thal_2	thal_3	ca_1	...	s
0	52	125	212	168	1.0	0	False	False	True	False	...	
1	53	140	203	155	3.1	0	False	False	True	False	...	
2	70	145	174	125	2.6	0	False	False	True	False	...	
3	61	148	203	161	0.0	0	False	False	True	True	...	
4	62	138	294	106	1.9	0	False	True	False	False	...	

5 rows x 23 columns

```
X = Dum_DF.drop(['target'], axis=1)
y = Dum_DF['target']
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train , y_test = train_test_split(X, y, test_size=0.2 , random_state=42)
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.fit_transform(x_test)
```

```
from sklearn.neural_network import MLPClassifier
from sklearn.datasets import make_classification
from sklearn.model_selection import GridSearchCV
```

```
hyper_parameters = {'batch_size':[ 'auto',100], 'max_iter':[200,500], 'hidden_layer_sizes':[5,(5,5,5)], 'learning_rate_init': [0.05,0.01,0.001]}
gs =GridSearchCV(MLPClassifier(),hyper_parameters,scoring='roc_auc',n_jobs=-1,return_train_score=False,verbose=0,cv=5)
clf =gs.fit(x_train, y_train)
print('The best combination is:')
print(clf.best_params_)

print('The best Accuracy is:')
print(clf.best_score_)
```

```
The best combination is:
{'batch_size': 'auto', 'hidden_layer_sizes': 5, 'learning_rate_init': 0.05, 'max_iter': 200}
The best Accuracy is:
0.9884341648051626
```

```
#Use random forest classifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

```
forest = RandomForestClassifier(n_estimators=10,criterion='entropy',random_state=1)
```

```
forest.fit(x_train, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(criterion='entropy', n_estimators=10, random_state=1)
```

```
model=forest
```

```
model.score(x_test,y_test)
```

```
0.9609756097560975
```

```
#XGB
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score
model = XGBClassifier(n_jobs=5,learning_rate=0.005)
model.fit(x_train,y_train)
```

```
▼ XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=0.005, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               multi_strategy=None, n_estimators=None, n_jobs=5,
               num_parallel_tree=None, random_state=None, ...)
```

```
predictions = model.predict(x_test)
```

```
model.score(x_test,y_test)
```

```
0.8292682926829268
```

```
import pandas as pd
```

```
# Dictionary containing model names and their accuracy
```

```
accuracy_data = {
    "Model": ["Random Forest", "MLP", "XG Boost"],
    "Accuracy (%)": [ 96.09, 98.84, 82.92]
}
```

```
# Create a DataFrame from the accuracy data
```

```
accuracy_df = pd.DataFrame(accuracy_data)
```

```
# Display the DataFrame
```

```
print(accuracy_df)
```

```
➡
```

	Model	Accuracy (%)
0	Random Forest	96.09
1	MLP	98.84
2	XG Boost	82.92